

# Class 07: Machine Learning 1


Olivia Baldwin

Before we get into clustering methods let's make some sample data to cluster where we know what it should look like.

To help, I will use the `rnorm()` function.

```
hist(c(rnorm(15000, mean=-3), rnorm(15000, mean=3)))
```

**histogram of `c(rnorm(15000, mean = -3), rnorm(15000, mean`**




`c(rnorm(15000, mean = -3), rnorm(15000, mean = 3))`

```
n=30
x <- c(rnorm(n, mean=3), rnorm(n, mean=-3))
y <- rev(x)
z <- cbind(x,y)
```

```
#cbind = by column
#rbind = by row

plot(z)
```



##K-means clustering The function in base R for k-means clustering is called `kmeans()`. - argument “x” = input and “centers” = how many clusters you want - “clustering vector” tells me which cluster each data point belongs to - “sum of squares” tells me how far each point in cluster are in the center/how tightly grouped (good) the cluster is

```
help(kmeans)
```

```
starting httpd help server ... done
```

```
km <- kmeans(z, centers=2)
km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	-2.585614	2.576812
2	2.576812	-2.585614

```
1 -2.852131  2.801046
2  2.801046 -2.852131
```

Clustering vector:

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Within cluster sum of squares by cluster:

```
[1] 43.83562 43.83562
(between_SS / total_SS =  91.6 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

Can get to the “components” by printing them out like you would columns (using a “\$”).

```
#examples of printing components
```

```
km$cluster
```


```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

```
km$centers
```

```
      x      y
1 -2.852131  2.801046
2  2.801046 -2.852131
```

Plot Z with the clusters colored and add cluster centers:

```
plot(z, col= km$cluster)
points(km$centers, col="blue", pch=15, cex=2)
```




```
#pch changed the point type to the squares  
#cex expanded the characters to size 2 instead of 1
```

Can you cluster our data in z into four clusters?

```
km4 <- kmeans(z, centers=4)
```

```
plot(z, col= km4$cluster)  
points(km4$centers, col="blue", pch=15, cex=2)
```



#this clustering has no basis, usually use a "scree plot" to look for an inflection point to

##Hierarchical Clustering

The main function for hierarchical clustering is `hclust()`. - can do top down or bottom up

Unlike `kmeans()` I cannot just pass in my data as an input, I first need a distance matrix from my data.

```
d <-dist(z)
hc <- hclust(d)
hc
```


Call:  
`hclust(d = d)`

Cluster method : complete  
Distance : euclidean  
Number of objects: 60

There is a specific `hclust plot()` method...

```
plot(hc)
abline(h=10, col="red")
```


## Cluster Dendrogram



```
#the line represents where you would "cut" the branches
#cut branches in the "easiest" spot, meaning the tallest height to determine clusters
```

To get my clustering result (i.e. the membership vector) I can “cut” the dendrogram at a given height. To do this I will use the `cutree()` function.

```
groups <- cutree(hc, h=10)
plot(z, col=groups)
```



## Principal Component Analysis

“Principal component analysis (PCA) is a well established” multivariate statistical technique” used to reduce the dimensionality of a complex data set to a more manageable number (typically 2D or 3D). This method is particularly useful for highlighting strong patterns and relationships in large datasets (i.e. revealing major similarities and differences) that are otherwise hard to visualize.”

### PCA for UK foods

```
url <- "https://tinyurl.com/UK-foods"  
uk <- read.csv(url)
```

Q1: 17 rows and 5 columns

```
dim(uk)
```

```
[1] 17  5
```

```
head(uk)
```

	X	England	Wales	Scotland	N.Ireland
1	Cheese	105	103	103	66
2	Carcass_meat	245	227	242	267
3	Other_meat	685	803	750	586
4	Fish	147	160	122	93
5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139

```
uk <- read.csv(url, row.names=1)
uk
```


	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494
Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41

Q2: I like the alternate option better to set the row names when you import so you don't have to override the variable.

Barplot:


```
barplot(as.matrix(uk), beside=T, col=rainbow(nrow(uk)))
```






Q3: Changing the beside argument will change the barplot.

```
barplot(as.matrix(uk), beside=FALSE, col=rainbow(nrow(uk)))
```



Q5: each y axis is aligned to the country in its row and being on the diagonal means the other country on each x matches the country on the y.

```
pairs(uk, col=rainbow(10), pch=16)
```



Q6: There are several points that are very different in the N.Ireland row (i.e. light blue, yellow, green, and orange that are higher on the x axis)

##PCA to the rescue

The main function to do PCA in base R is `prcomp()`.

Note that I need to take the transpose of this particular data as that is what the `prcomp()` help page asked for.

```
pca <- prcomp( t(uk) )
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	3.176e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Lets see what is inside our result object `pca` that we just calculated.

```
attributes(pca)
```

```
$names
[1] "sdev"      "rotation" "center"    "scale"     "x"


$class
[1] "prcomp"
```

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-4.894696e-14
Wales	-240.52915	-224.646925	-56.475555	5.700024e-13
Scotland	-91.86934	286.081786	-44.415495	-7.460785e-13
N.Ireland	477.39164	-58.901862	-4.877895	2.321303e-13

To make our main result figure, called a “PC plot” (or score plot or ordination plot or PC1 vs. PC2).

```
plot(pca$x[,1], pca$x[,2], col=c("black", "red", "blue", "green"), pch=16, xlab="PC1 (67.4%)
```



```
#this will access the first and second column of the PC table output (a.k.a. PC1 and PC2)
```

## Variable Loadings Plot

This plot can show how these countries differ from each other. The positive means it is moving towards the positive end of PCA1.

```
## Lets focus on PC1 as it accounts for > 90% of variance  
par(mar=c(10, 3, 0.35, 0))  
barplot( pca$rotation[,1], las=2 )
```

